Message Queues: RabbitMQ

To create a simple messaging system, I used Producer and Consumer applications, paired with RabbitMQ. The producer loops a million times to post a message to the queue, and the consumer consumes the messages from the queue. The code for this is available on Github.

After installing and starting RabbitMQ on my system, I started both applications together so the consumer started reading from the queue in real-time. After the producer application had sent a million messages and consumer had received them, I was able to confirm the same from the logs by maintaining a global variable 'total count' in the consumer application. It was observed that all 1,000,000 messages were consumed correctly.

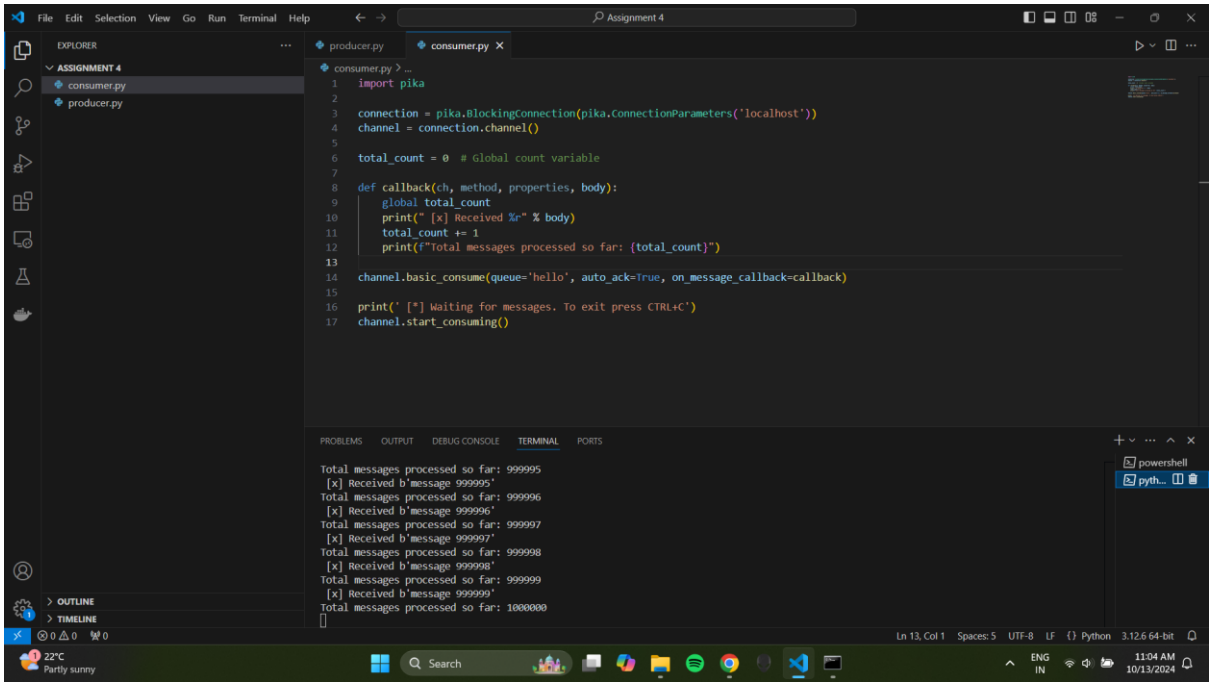

Having enabled the RabbitMQ Management Plugin, I was able to view some additional metadata on the 15672 port on localhost. It was observed that the consumption of the messages was entirely real-time as it showed no lags on the graphs. Through these two, we can confidently claim that not a single message was dropped in the queue and they were successfully consumed.